

Designing Data-Intensive Applications

Notes

Contents

How to Use This Template	2
1 Chapter # – Title	3
1.1 General	3

How to Use This Template

- Create a new **Chapter** section for each audiobook chapter.
- Capture key ideas, notable quotes, and takeaways in the provided subsections.
- Add action items or open questions you want to revisit.

1 Chapter # – Title

1.1 General

- Redis, Apache Kafka (boundaries blurred)
- Single tool can't any longer meet data needs
- work is broken down into tasks into multiple tools
- memcached, elastic search, solar – application code keeps all this synced
- API hides implementation details from clients
 - how do you ensure data remains correct and complete?
 - how do you scale?
 - what API is best
 - how do you have an issue but the client can still be fine
- Book considers 3 concerns: reliability (toyota style), scalability, maintainability (changing people should be able to work on the system productively)
 - following chapters: architectures and algos important to achieve these
 - Reliability:
 - faults and fault tolerant systems in reliability
 - fault is one component of a system deviating from its spec
 - failure: system as a whole stops providing the service to the user
 - fault tolerance mechanisms are important
 - Netflix Chaos Monkey: a revolutionary tool that randomly disables virtual machines and services in its production environment to test system resilience

- hard disk MTTF (mean time to failure) is 10-50 years, 10000 disks, on average 1 disk to die per day
- mo machines mo problems
- Virutal Machine instances stop running without warning because it is sometimes needed somewhere else
- Software Errors:
 - independent events generally, a large number don't fail all at once
 - automated testing: config changes -> possible issues -> slow rollout so only a few users are affected if at all
 - telemetry: track and monitor failures
 - discussing scalability is trying to figure out how to cope with growth in whichever way you are growing (current load, and then incr. load)
 - twitter example (Nov 2012 data): 4600 req./s, views 300,000 req/s, tweet volume + fan out. Two way s of implementing follower followee. Global collection is an option and then you merge those where follower and followee (original approach). do more work at write time and less at read time so just write the tweets on their home timelines (new approach). "Fan out load". Twitter now hybrids b/w the 2. Celebs are fetched separately like in Approach 1.
 -

Key Ideas

-
-
-

Architecture / Systems Mentioned

-

Concepts and Definitions

- **Term:** Definition or explanation.

Diagrams / Mental Models

-

Notable Quotes

- “”

Questions / Confusions

-

Practical Takeaways

-

Action Items

-