

Data Science Notes

Bhavjot Khurana

November 5, 2025

Contents

1	Introduction	3
2	Topic 1: Linear Regression	3
2.1	Definition	3
2.2	Key Formulas	3
2.3	Error Function	3
2.4	Mean Squared Error (MSE)	4
2.5	How the Line is Fitted	4
2.6	Hypothesis Testing and p-values	4
2.7	Residual Standard Error (RSS and TSS)	4
2.8	Interpretations	5
2.9	Example: Simple Linear Regression Output	5
3	Topic 2: Multiple Linear Regression	5
3.1	Definition	5
3.2	F-statistic and Interpretation	6
3.3	Example: Multiple Linear Regression Output	6

4 Topic 3: Machine Learning	6
4.1 Supervised Learning	6
4.2 Unsupervised Learning	7
5 Conclusion	7

1 Introduction

Using the freecodecamp.org video as a reference (link: <https://www.youtube.com/watch?v=XU5pw3QR>). This document contains notes on data science topics covered in the video.

2 Topic 1: Linear Regression

2.1 Definition

Linear regression models the expected value of a response variable y as a linear function of a single predictor x :

$$y_i = \beta_0 + \beta_1 x_i + \varepsilon_i, \quad \varepsilon_i \sim \text{Normal}(0, \sigma^2)$$

It assumes the errors ε_i are independent, normally distributed, and have constant variance.

2.2 Key Formulas

Ordinary least squares (OLS) estimates the slope and intercept by minimizing the sum of squared residuals:

$$\begin{aligned}\hat{y}_i &= \hat{\beta}_0 + \hat{\beta}_1 x_i \\ \hat{\beta}_1 &= \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}, \quad \hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}\end{aligned}$$

2.3 Error Function

The residual for observation i captures the prediction error:

$$e_i = y_i - \hat{y}_i$$

Residuals tell us how far each point lies from the regression line; examining their pattern helps spot outliers or violations of model assumptions.

2.4 Mean Squared Error (MSE)

Mean Squared Error is the average of squared residuals:

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Errors are squared to penalize large deviations more heavily, eliminate sign cancellation, and produce a smooth, differentiable loss function that calculus-based solvers can optimize.

2.5 How the Line is Fitted

Fitting the line involves solving the optimization problem:

$$\min_{\beta_0, \beta_1} \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2$$

Setting the partial derivatives to zero yields the normal equations (matrix form: $(X^\top X)\hat{\beta} = X^\top y$). Software uses closed-form solutions for small problems or matrix decompositions (QR or SVD) for numerical stability. Gradient methods (e.g., gradient descent) offer scalable alternatives for very large datasets.

2.6 Hypothesis Testing and p-values

To test whether x helps explain y , use the t-test for the slope. The null hypothesis $H_0 : \beta_1 = 0$ indicates no linear relationship. The statistic

$$t = \frac{\hat{\beta}_1}{\text{SE}(\hat{\beta}_1)}$$

follows a t-distribution with $n - 2$ degrees of freedom. A small p-value suggests the predictor provides statistically significant explanatory power.

2.7 Residual Standard Error (RSS and TSS)

Two sums of squares quantify variability:

$$\text{RSS} = \sum_{i=1}^n (y_i - \hat{y}_i)^2, \quad \text{TSS} = \sum_{i=1}^n (y_i - \bar{y})^2$$

Residual Standard Error (RSE) estimates the typical size of residuals in the units of y :

$$\text{RSE} = \sqrt{\frac{\text{RSS}}{n - 2}}$$

Lower RSE indicates tighter fit around the regression line.

2.8 Interpretations

- **Slope** $\hat{\beta}_1$: Expected change in y for a one-unit increase in x .
- **Intercept** $\hat{\beta}_0$: Expected value of y when $x = 0$, useful when $x = 0$ is meaningful.
- $R^2 = 1 - \frac{\text{RSS}}{\text{TSS}}$: Proportion of variance in y explained by the model.
- **Practical use**: Combine coefficient estimates with confidence or prediction intervals to communicate both central tendency and uncertainty in predictions.

2.9 Example: Simple Linear Regression Output

Using the `Advertising` dataset (TV spend predicting sales), the `statsmodels` summary reports the following key statistics:

Statistic	Estimate	Interpretation
Intercept ($\hat{\beta}_0$)	7.03	Baseline sales when TV spend is \$0k.
TV coefficient ($\hat{\beta}_1$)	0.0475	Each \$1k on TV adds approximately 0.048k units sold.
R^2	0.612	61.2% of variance in sales explained by TV.
F-statistic	312.1 ($p < 10^{-40}$)	The model is highly significant overall.

The printed OLS table also shows small standard errors and very large t-statistics for the slope, so the p-value for the TV coefficient is essentially zero. Residual diagnostics (Omnibus, Jarque-Bera, Durbin-Watson) indicate roughly normal residuals and limited autocorrelation for this example.

3 Topic 2: Multiple Linear Regression

3.1 Definition

Multiple Linear Regression (MLR) extends the linear model to p predictors:

$$y_i = \beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip} + \varepsilon_i$$

Coefficients measure the expected change in y for a one-unit change in the corresponding predictor while holding others constant.

3.2 F-statistic and Interpretation

The model-wide F-test compares the explained variance to unexplained variance:

$$F = \frac{(TSS - RSS)/p}{RSS/(n - p - 1)}$$

A large F-statistic (with a small p-value) indicates that, taken together, the predictors explain significantly more variation than an intercept-only model. If the F-test is not significant, it implies the collective set of predictors may not offer meaningful predictive power beyond the mean of y .

3.3 Example: Multiple Linear Regression Output

Fitting TV, radio, and newspaper spend simultaneously to predict sales yields the following summary highlights:

Statistic	Estimate	Interpretation
Intercept	2.94	Baseline sales when all spends are zero.
TV coefficient	0.0458	Positive, highly significant effect ($p < 10^{-50}$).
Radio coefficient	0.1885	Strong positive effect ($p < 10^{-40}$).
Newspaper coefficient	-0.0010	Not significant ($p \approx 0.86$).
R^2	0.897	89.7% of sales variance explained jointly.
F-statistic	570.3 ($p < 10^{-90}$)	Model is overwhelmingly significant.

Compared with the simple-regression model, R^2 jumps from 0.61 to 0.90, showing the value of additional predictors. However, the newspaper coefficient's large p-value signals it may not contribute meaningfully, suggesting a model refinement (e.g., dropping the variable) could be appropriate.

4 Topic 3: Machine Learning

4.1 Supervised Learning

Explain supervised learning here.

4.2 Unsupervised Learning

Explain unsupervised learning here.

5 Conclusion

Summarize your notes here.